




Chapter 4, Part 5: Linear System Solvers

Xiaofeng Liu, Ph.D., P.E.
 Assistant Professor
 Department of Civil and Environmental Engineering
 Pennsylvania State University
 xliu@engr.psu.edu



Linear system solvers

Linear System of Equations

3/31

This chapter is not intended to be a complete lecture on linear system solvers. For these, you need to take a course or read some good books.

- ▶ Yousef Saad, Iterative Methods for Sparse Linear Systems, second edition
http://www-users.cs.umn.edu/~saad/IterMethBook_2ndEd.pdf

The content of this lecture is mostly from online forum, open reports, and other people's lecture notes.

Linear System of Equations

4/31

In OpenFOAM[®], the major part related to the linear system of equations is in `src/OpenFOAM/matrices/lduMatrix/`

It has the following sub-directories:

- ▶ `lduAddressing`: addressing for lower triangle, diagonal, and upper triangle coefficients of the matrix
- ▶ `lduMatrix`: the ldu matrix itself
- ▶ `preconditioners`: preconditioners for linear solvers
- ▶ `smoothers`: smoothers for linear solvers
- ▶ `solvers`: linear solvers

Nomenclature

- ▶ For each computational cell, we will create an equation

$$a_P \phi_P + \sum_N a_N \phi_N = b$$

- ▶ Equations form a **linear system** or a matrix

$$A\phi = b$$

where A contain matrix coefficients, ϕ is the value of ϕ_P in all cells and b is the right-hand-side

- ▶ A is potentially very big: $M \times M$, M being the number cells
- ▶ A is a **square matrix**: the number of equations equals the number of unknowns
- ▶ A is sparse, i.e., very few coefficients are non-zero. The reason is discretization stencil is local.
- ▶ A needs good format to reduce storage size and computational efficiency

Since most of CFD codes have to deal with sparse matrix, we will only look at sparse cases.

- ▶ Sparse matrix format: Only non-zero coefficients will be stored.
 - Considerable savings in memory
 - Need a mechanism to indicate the position of non-zero coefficients
 - The format is static, which imposes limitations on the operations: if a coefficient is originally zero and not stored, it is very expensive to set its value. This is usually termed a zero fill-in condition.
 - There are several sparse matrix storage formats.
 - coordinate format
 - compressed row storage (CRS) format
 - compressed column storage (CCS) format
 - ...

More information:

A survey of sparse matrix storage formats:

http://netlib.org/linalg/html_templates/node90.html

Section 3.4 of Saad (2003) (Free online book).

Design of sparse matrix format needs to consider:

- ▶ Storage overhead
- ▶ Access of matrix elements
- ▶ Operations on the matrix (e.g., LU decomposition)
- ▶ Operations with the matrix (e.g., matrix-vector multiplication)

Example: Compressed Row Storage (CRS) Format

- ▶ Three arrays: *value* for the matrix coefficients, *column* and *row* for addressing
- ▶ Coefficients ordered row-by-row
- ▶ The *column* array records the column index for each coefficients. Size of *column* array equal to the number of off-diagonal coefficients
- ▶ The *row* array records the start and end of each row in the *column* array. Thus, row i has got coefficients from $row[i]$ to $row[i + 1]$ (not included). Size of *row* arrays equal to number of rows + 1

$$A = \begin{bmatrix} 1 & & & & \\ & 2 & 5 & & \\ & 3 & 6 & & 9 \\ & 4 & & 8 & \\ & & 7 & & 10 \end{bmatrix} \quad (1)$$

$$\begin{aligned} val &= [1, 2, 5, 3, 6, 9, 4, 8, 7, 10] \\ column &= [1, 2, 3, 2, 3, 5, 2, 4, 3, 5] \\ row &= [1, 2, 4, 7, 9, 11] \end{aligned}$$

Example: Arrow Format (used in OpenFOAM®)

- ▶ Actual data stored in the `lduMatrix` class and addressing implemented in the `lduAddressing` class
- ▶ Arbitrary sparse format. Diagonal coefficients typically stored separately
- ▶ Coefficients in 2-3 arrays: diagonal, upper and lower triangle in `lduMatrix` class, which has member functions:

```
const scalarField & lower () const  
const scalarField & diag () const  
const scalarField & upper () const
```

- ▶ Diagonal addressing implied (cells)
- ▶ Off-diagonal addressing in 2 arrays: “owner” (row index) “neighbour” (column index) array.
- ▶ Size of lower and upper addressing = the number of coefficients = number of internal faces
- ▶ If the matrix coefficients are symmetric, only the upper triangle is stored – a symmetric matrix is easily recognised and stored only half of coefficients

Matrix Storage Formats

10/31

For example, given a simple 3×3 orthogonal mesh (in 2-dimensions):

-1-	-2-	-3-
-6-	-5-	-4-
-7-	-8-	-9-

The coefficient matrix for diffusion equation might look like:

i	-1-	-2-	-3-	-4-	-5-	-6-	-7-	-8-	-9-
-1-	X	N				N			
-2-	O	X	N		N				
-3-		O	X	N					
-4-			O	X	N				N
-5-		O		O	X	N		N	
-6-	O				O	X	N		
-7-						O	X	N	
-8-					O		O	X	N
-9-				O				O	X

Where X indicates diagonal cells, O and N indicate non-zero off-diagonal cells.

- ▶ Since the matrix is sparse, only the non-zero entries need to be stored.
 - All diagonal coefficients a_p are non-zero.
 - Sparsity only refers to the off-diagonal.
 - lower and upper triangles have $N(N - 1)$ entries; most of them are zero.
 - The non-zero coefficients correspond to “internal faces” in the mesh: one appears in the lower triangle and one in upper triangle
 - The two coefficients corresponding to one internal face could be the same (e.g., diffusion term) or different (e.g., advection term)
 - So Laplace (Poisson) equation results a symmetric matrix A , while others (advection-diffusion) results in asymmetry.




Figure: An example of sparse matrix (source:wikipedia)

- ▶ The storage of matrix is implemented in *lduAddressing* class in OpenFOAM®.
- ▶ It works closely with the mesh data structure
 - points: (x, y, z) coordinates
 - faces: list of points whose order makes sure the face normal pointing from owner to neighbor (using right hand rule). Boundary faces always point outward.
 - face owner and neighbor: cells on both sides of the face. The owner always has lower number. Boundary faces only have owner and no neighbor.
 - The `owner` and `neighbor` files inside `constant/polyMesh` directory: `owner` has all the faces (internal + boundary); `neighbor` only have internal faces and thus is shorter and `owner` by the number of boundary faces.
 - Thus `neighbor`, i.e., internal faces, determines the off-diagonal coefficients.
 - the
 - Boundary patches are defined in `boundary` file: name, type, number of faces `nFaces`, start index `startFace` in the face list.
 - Remember the effect of boundary faces: only affect the diagonal and the right hand side of the linear equation. The class `fvMatrix` (inherited from `lduMatrix` has two functions:
`addBoundaryDiag();`
`addBoundarySource();`

- ▶ *lduAddressing* works by recognizing that:
 - Every cell has a diagonal coefficient; and
 - Every cell has off-diagonal coefficients for each of their neighbors.
- ▶ Therefore the diagonal coefficients are stored in an “N”-long vector array, where “N” is the number of cells.
- ▶ The number of off-diagonal coefficients is equal to the number of cell-pairs (internal faces) that directly influence one another in the linearized equations.
- ▶ At the matrix level in OpenFOAM[®], this only includes “adjacent” cells. Therefore the number of off-diagonal coefficients is equal to the number of shared “faces” (internal faces) in the mesh.

1duAddressing:

- ▶ The diagonal coefficients are indexed by "cell index". So `Diag[cellI]` is the diagonal coefficient for `cellI`.
- ▶ The off-diagonal coefficients (lower and upper triangles) are indexed by "face index".
- ▶ For example, to recover the full 2D matrix:

```
//A defined as a 2D matrix
```

```
//Loop over all the cells to assemble the diagonal coefficients  
for (label cellI=0; cellI<Diag.size(); cellI++)
```

```
{
```

```
    A[cellI, cellI] = Diag[cellI];
```

```
}
```

```
//Loop over all the internal faces to assemble the  
//off-diagonal coefficients.
```

```
for (label face=0; face<l.size(); face++)
```

```
{
```

```
    A[l[face],u[face]] = lower[face];
```

```
    A[u[face],l[face]] = upper[face];
```

```
}
```

lduAddressing:

- ▶ l and u are lower and upper triangle addressing
- ▶ The lower triangle is called “owner” coefficients which comes from the face’s owner cell
- ▶ The upper triangle is called “neighbor” coefficients which comes from the face’s neighbor cell

FVM and Matrix Structure: `fvMatrix`

- ▶ A class inherited from `lduMatrix`
- ▶ It adds:
 - Reference to the solution variable ϕ
 - Dimensions: M, L, T , etc.
 - Stores the right hand side b .
 - Handles boundary condition effects in the discretization:
`addBoundaryDiag()`
`addBoundarySource()`
 - ...

The Role of a Linear Solver

- ▶ Numerical simulation software will spend a big portion of CPU time solving $Ax = b$:
- ▶ Performance of linear solvers is absolutely critical for the performance of the simulation
- ▶ The performance depends on several factors:
 - the linear solver itself
 - the property of matrix A due to discretization schemes

Category of linear solvers:

► **Direct solver**

- Example: LU decomposition
- Not attractive to solve large linear system
- Computational time grows very quickly as $O(N^3)$ or $O(N^2)$
- Exact solution is not necessary since the linear system is an approximation
- Breaks sparsity

► **Iterative solver:**

- The algorithm will start from an initial solution and perform a number of operations which will result in an improved solution.

$$\phi^{(0)} \rightarrow \phi^{(1)} \rightarrow \phi^{(2)} \dots \rightarrow \phi^{(i)} \rightarrow \dots \rightarrow \phi$$

- Usually takes advantage and preserve the sparsity pattern
- Types:
 - Relaxation methods: Jacobi-, Gauss-Seidel-Relaxation, ...
 - Krylov subspace solvers: PCG, PBiCG, etc.
 - Multigrid methods

- ▶ Performance of iterative solvers depends on the matrix characteristics.

$$Ax = b$$

- ▶ The solver operates by incrementally improving the solution, i.e., reducing the error
- ▶ If the error is amplified in the iterative process, the solver diverges
- ▶ Preconditioners: make sure the convergence of the preconditioned system is much faster than the original one.

$$M^{-1}Ax = M^{-1}b$$

- ▶ Smoothers: smoothing algorithms for multi-grid method guarantee that the approximate solution after each solver iteration will be closer to the exact solution than all previous approximation. An example of a smoother is the Gauss-Seidel algorithm

Matrix Characterisation

- ▶ A matrix is **sparse** if it contains only a few non-zero elements
- ▶ A sparse matrix is **banded** if its non-zero coefficients are grouped in a stripe around the diagonal
- ▶ A sparse matrix has a **multi-diagonal structure** if its non-zero off-diagonal coefficients form a regular diagonal pattern
- ▶ A **symmetric** matrix is equal to its transpose

$$[A] = [A]^T$$

- ▶ A matrix is **positive definite** if for every $[\phi] \neq [0]$

$$[\phi]^T [A] [\phi] > 0$$

Matrix Characterisation

- ▶ A matrix is **diagonally dominant** if in each row the sum of off-diagonal coefficient magnitudes is equal or smaller than the diagonal coefficient

$$a_{ii} \geq \sum_{j=1}^N |a_{ij}| \quad ; \quad j \neq i$$

and for at least one i

$$a_{ii} > \sum_{j=1}^N |a_{ij}| \quad ; \quad j \neq i$$

- ▶ Matrix form of the system we are trying to solve is

$$[A][\phi] = [b]$$

- ▶ The exact solution can be obtained by inverting the matrix $[A]$:

$$[\phi] = [A]^{-1} [b]$$

- ▶ Iterative solvers start from an approximate solution ϕ^0 and generates a sequence of solution estimates ϕ^k , where k is the iteration counter
- ▶ convergence is measured through residual:

$$[r]^k = [b] - [A][\phi]^k$$

Residual is a vector showing how far is the current estimate $[\phi]^k$ from the exact solution $[\phi]$.

- ▶ How to measure residual? Norm of $||r||$

$$||r|| = \sum_{j=1}^N |r_j|$$

which is the absolute error norm.

- ▶ We usually need to normalize it for easy quantification (compare with tolerance in `fvSolution`).

$$R_s = \frac{||r||}{normFactor} < tolerance?$$

Definition of a Residual

24/31

- ▶ How to define the normFactor? OpenFOAM® does it in a slightly unusual way.

Residual is defined as

$$r = b - A\phi$$

OpenFOAM® defines a reference solution value ϕ_{ref} = volume average of ϕ over the whole domain, then

$$w_A = A\phi^k$$

$$p_A = A\phi_{ref}^k$$

The scaling factor is then calculated as

$$scaleFactor = \sum |w_A - p_A| + |b - p_A| + 10^{-20}$$

The scaled absolute error norm:

$$r_s = \frac{\sum |b - w_A|}{scaleFactor}$$

- ▶ This definition of `normFactor` breaks down when the solution $\phi \equiv 0$ and `normFactor` is a very small number.

- ▶ The code implementation is in class `solverPerformance` and the linear system solver such as PCG

```
src/OpenFOAM/matrices/lduMatrix/solvers/PCG
```

- ▶ You can print out the value of `normFactor` by setting the debug level to ≤ 2 because all solvers have:

```
if (lduMatrix::debug >= 2)
{
    Info<< "    Normalisation factor = " << normFactor << endl
}
```

- ▶ The actual calculation of `normFactor` is in the `lduMatrix` class:

```
src/OpenFOAM/matrices/lduMatrix/lduMatrix
```

```
...
```

```
Foam::scalar Foam::lduMatrix::solver::normFactor(...)
```

- ▶ To check the convergence, one can also use relative convergence measure (compare with `relTol` in `fvSolution`)

$$\frac{\|r_k\|}{\|r_0\|} < relTol ?$$

- ▶ If both tolerance and `relTol` are specified, the stopping criterion is “OR” relationship, . This is coded in

`/src/OpenFOAM/matrices/LduMatrix/LduMatrix`

```
if
(
    finalResidual_ < Tolerance
    || (
        RelTolerance
        > small_*pTraits<Type>::one
        && finalResidual_ < cmptMultiply(RelTolerance, initial
    )
)
{ converged_ = true;}
```

Examples of Iterative Solvers

27/31

In OpenFOAM[®], the linear solvers are in

`src/OpenFOAM/matrices/lduMatrix/solvers`

- ▶ BICCG: Diagonal incomplete LU preconditioned BiCG solver
- ▶ diagonalSolver: diagonal solver for both symmetric and asymmetric problems
- ▶ GAMG: Geometric agglomerated algebraic multigrid solver (also named Generalised geometric-algebraic multi-grid in the manual)
- ▶ ICCG: Incomplete Cholesky preconditioned Conjugate Gradients solver
- ▶ PBiCG: Preconditioned bi-conjugate gradient solver for asymmetric lduMatrices using a run-time selectable preconditioner
- ▶ PCG: Preconditioned conjugate gradient solver for symmetric lduMatrices using a run-time selectable preconditioner
- ▶ smoothSolver: Iterative solver using smoother for symmetric and asymmetric matrices which uses a run-time selected smoother

In OpenFOAM[®], the preconditioners are in

`src/OpenFOAM/matrices/lduMatrix/preconditioners`

- ▶ `diagonalPreconditioner`: Diagonal preconditioner for both symmetric and asymmetric matrices.
- ▶ `DICPreconditioner`: Simplified diagonal-based incomplete Cholesky preconditioner for symmetric matrices (symmetric equivalent of DILU).
- ▶ `DILUPreconditioner`: Simplified diagonal-based incomplete LU preconditioner for asymmetric matrices.
- ▶ `FDICPreconditioner`: Faster version of the `DICPreconditioner` diagonal-based incomplete Cholesky preconditioner for symmetric matrices (symmetric equivalent of DILU)
- ▶ `GAMGPreconditioner`: Geometric agglomerated algebraic multigrid preconditioner
- ▶ `noPreconditioner`: Null preconditioner for both symmetric and asymmetric matrices.

In OpenFOAM[®] , the smoothers are in

`src/OpenFOAM/matrices/lduMatrix/smoothers`

- ▶ DIC: diagonal-based incomplete Cholesky smoother for symmetric matrices.
- ▶ DICGaussSeidel: Combined DIC/GaussSeidel smoother for symmetric matrices
- ▶ DILU: diagonal-based incomplete LU smoother for asymmetric matrices.
- ▶ DILUGaussSeidel: Combined DILU/GaussSeidel smoother for asymmetric matrices
- ▶ GaussSeidel

Examples of Iterative Solvers

30/31

In OpenFOAM[®], the specification of linear solver is in file `fvSolution`:

```
solvers
{
    p PCG
    {
        preconditioner DIC;
        tolerance 1e-06;
        relTol 0;
    };

    U PBiCG
    {
        preconditioner DILU;
        tolerance 1e-05;
        relTol 0;
    };
}
```

Examples of Iterative Solvers

31/31

Demonstration of the effects of different choices of linear solvers.

- ▶ Pressure linear solver PCG with different preconditioners
- ▶ Different pressure linear solvers